

## *In-class Programming Activity 12*

### *Math 253: Statistical Computing & Machine Learning*

#### *Testing LDA and QDA*

Today, you're going to generate simulated data to test out LDA and QDA.

You'll make two data sets, each of which will have two classes of outcomes: "red", and "blue"

- SameSigma where cases in the two classes have the same covariance matrix for the predictor variables.
- DifferentSigma where the two classes have different covariance matrices.

#### *Generating simulated data*

- Create an object `n_cases` with the value 100.
- Create an object `red_mean` with the value `c(1, 0)`.
- Create an object `green_mean` with value `c(0, -1)`.
- Create an object `blue_mean` with value `c(-1, 1)`.
- Create an object `covar_1` which is a  $2 \times 2$  matrix with 3 and 1 on the diagonal and  $-1.7$  on the off-diagonal. It should be structured as a covariance matrix.
- Similarly, create an object `covar_2`, a covariance matrix with 2 and 3 on the diagonal and 1.5 on the off-diagonal.
- Create three matrices, one, two, and three. Each should have the appropriate shape for data from `n_cases` cases and two variables. The variables should each be `n_cases` random draws from a  $N(0, 1^2)$  distribution, that is, a normal distribution with mean zero and standard deviation 1.
- Create a matrix `red` that is one times the Cholesky decomposition of `covar_1`. The `red` matrix will contain correlated random variables with a covariance of approximately `covar_1`.
- Similarly, create a matrix `green` which will be two times the Cholesky decomposition of `covar_1`.
- Also create a matrix `blue` which will be three times the Cholesky decomposition of the other covariance matrix, `covar_2`.
- Modify the `red`, `green` and `blue` matrices by adding to each column a value for the mean drawn from `red_mean`, `green_mean`, and `blue_mean` respectively. That is, for `red`, add 1 to the first column and 0 to the second.

In specifying the normal distribution, one needs to decide whether to report the standard deviation or the variance. R uses `sd=`. To help to eliminate ambiguity in mathematical notation, the form  $1^2$  is used simply as a reminder that the quantity is a variance.

*Hint:* It won't work to do the obvious, simple thing, e.g. add `red_mean` to `red`. There are many ways to construct a statement that works. Among others, there's a way using `outer()`, a way using `matrix()`, and even a way using `t()` twice.

- Create three data frames, each with variables  $x$ ,  $y$ , and  $class$ .
  - Red will have  $x$  as the first column of `red`,  $y$  as the second column of `red`, and `class` set equal to the string “red”.
  - Blue will be the same thing but using the columns of `blue` and the `class` set to “blue”
  - Green is similar, using the columns of `green` and the `class` “green”.
- Last step in generating the simulated data: make two data frames each of which combines “data” from two classes.

```
Sim_one <- rbind(Red, Green)
```

```
Sim_two <- rbind(Red, Blue)
```

### LDA and QDA

Fit a linear discriminant model `class ~ x + y` to the data in `Sim_one`. Call the model `mod_LDA_one`.

```
mod_LDA_one <- MASS::lda(class ~ x + y, data = Sim_one)
```

Then use the model to test the model on the same training data to which it was fit. Store the result in `test_LDA_one`.

```
test_LDA_one <- predict(mod_LDA_one, newdata = Sim_one)
```

The resulting object, `test_LDA_one`, is a list of three items. Make sure you understand what each of them is.

QDA works in the same way: the function is `qda()`.

### Confusion matrices

The confusion matrix compares the actual class to the predicted class from the model. It's straightforward to compute:

```
table(Sim_one$class, test_LDA_one$class)
```

- Compare the confusion matrix from LDA on `Sim_one` to that from QDA on `Sim_one`.<sup>1</sup> Which one shows better performance?
- Fit both LDA and QDA models to `Sim_two`. Which one performs better?

### Bigger $n$

The difference in performance of LDA and QDA in these examples is not so large that it's evident in a sample with 100 cases of each class. Go back and set `n_cases` to be 10000, and re-evaluate the confusion matrices.

### Above and beyond

Calculate the log likelihood for `mod_LDA_one` against the observations `Sim_one$class`.

You'll use `data.frame()` to construct the data frames. Make sure to give the optional argument `stringsAsFactors = FALSE`. This will let the `class` be stored as straightforward character strings that can be used in plotting to specify the color.




Figure 1: The two classes of cases, red and blue, in `Sim_two`.

<sup>1</sup> Use the names `mod_QDA_one` and `test_QDA_one` to store the fitted model and the test results from `predict()` respectively.